

1. Dataset Validation

The two result files were structurally consistent:

Validation item	Result
Rows in each dataset	1,000
Columns in each dataset	33
LLM metrics per example	9
Source records matched	1,000/1,000
Reward-function scores matched	100%
Evaluation errors	0
Maximum LLM total	45

The source hate speech, generated counter-narrative, reference counter-narrative, target group, dataset version, judge model, and provider were aligned row by row. Therefore, the observed differences originated from the LLM-judge scoring process rather than from differences in the input data.

2. Average Metric Scores

Metric	Integer scoring	Decimal scoring	Difference
PRS	3.364	3.315	-0.049
QS	3.231	3.278	+0.047
SAFE	3.722	3.836	+0.114
EMP	2.597	2.792	+0.195
CON	3.027	3.137	+0.110
PERS	3.027	3.137	+0.110
SPEC	3.437	3.502	+0.065

FLU	4.256	3.921	-0.335
TOX	3.951	3.757	-0.194
Mean total	30.612/45	30.674/45	+0.062
Mean percentage	68.03%	68.16%	+0.13 pp

At the aggregate level, the two procedures produced nearly identical overall scores. The decimal procedure increased the mean total by only 0.062 points out of 45.

However, individual metrics changed differently. The largest increase occurred for empathy, while the largest reductions occurred for fluency and non-toxicity.

3. Row-Level Agreement

Although the overall averages were similar, substantial differences were observed at the individual-example level.

Comparison measure	Result
Rows with at least one changed metric	938/1,000
Completely unchanged rows	62/1,000
Changed metric values	5,049/9,000
Percentage of changed metric values	56.1%
Rows with different total scores	888/1,000
Mean absolute total-score difference	5.18/45
Rows differing by more than 10 points	133
Rows differing by more than 20 points	18
Largest total-score increase	+33.0
Largest total-score decrease	-31.5

These findings show that positive and negative changes largely cancelled each other at the aggregate level. Therefore, similar dataset averages should not be interpreted as evidence of strong row-level agreement.

4. Decimal-Score Utilization

The decimal procedure did not consistently use the additional scoring precision.

Decimal compliance measure	Result
Rows containing at least one fractional score	161/1,000
Fractional metric values	931/9,000
Percentage of fractional metric values	10.34%
Batches containing fractional scores	26/125

Most scores in the decimal dataset remained whole-number values represented as decimals, such as 3.0, 4.0, and 5.0. Only 10.34% of all metric values used genuine fractional precision.

The concentration of fractional scores in 26 of 125 batches indicates that compliance with the decimal-scoring instruction varied across batches.

5. Examples of Large Differences

Row index	Integer total	Decimal total	Difference
23	6.0	39.0	+33.0
819	6.0	38.0	+32.0
406	45.0	13.5	-31.5
991	44.0	13.0	-31.0
783	43.0	13.0	-30.0

These extreme reversals are larger than what would normally be expected from changing only the numeric precision. They indicate variability in the judge's interpretation or application of the evaluation rubric.

6. Interpretation

The integer and decimal scoring procedures produced similar aggregate results:

- Integer mean: 30.612 out of 45
- Decimal mean: 30.674 out of 45

However, the row-level analysis revealed limited agreement. More than 93% of examples had at least one changed metric, and 56.1% of all metric values differed between the two runs.

Therefore, the decimal results should not be interpreted as simple refinements of the original integer scores. They represent a separate LLM evaluation run in which the model reconsidered many judgments rather than merely assigning more precise values around the original integer scores.

The decimal format provides greater theoretical precision, but the current run used that precision inconsistently. The principal concern is judge stability and rubric consistency rather than an error in the score calculations.

7. Conclusion

Both datasets are structurally valid, and their total and percentage calculations are correct. The overall evaluation outcome remained stable at the dataset level, with an average difference of only 0.062 points out of 45.

Nevertheless, substantial row-level discrepancies were found. The decimal run changed 56.1% of metric values and produced several extreme score reversals. Furthermore, only 10.34% of decimal metric values contained genuine fractional precision.

For future experiments, the decimal scoring procedure may be retained as the official method because it permits more fine-grained evaluation. However, results should be generated using a fixed prompt, model version, temperature, seed where supported, and consistent batch size. Multiple judging runs or agreement analysis are recommended before treating decimal scoring as more reliable than integer scoring.